

AI Champions Phase 1 - Appendix Q15

Risk register derived from the scored risk section of the application.

Risk	Likelihood	Impact	Mitigation	Residual
LLM accuracy on complex polypharmacy cases stays below target	High	High	Curriculum learning, expanded training data, ensemble fallback, safety layer catches unsafe output	Medium
Safety rule coverage misses rare combinations	Medium	High	DrugBank and guideline cross-check, pharmacist review, uncertain cases escalated to physician	Low
Diraya retrospective data quality issues reduce validation reliability	Medium	Medium	Month 1 data audit, cleaning, exclusion thresholds, large sample size preserves statistical power	Low
Health-system adoption remains slow after technical validation	Medium	Medium	Generate trust-focused evidence, co-design with practising clinician, de-risk in Phase 2 pilots	Medium
Key-person dependency on Sami and Dr. Trigos	Medium	Medium	Recruit ML Engineer, document validation methodology, use independent clinical reviewer support	Medium

Critical inputs

- Retrospective Diraya data secured through Dr. Trigos
- Commercially available GPU infrastructure
- Open-weight foundation models such as Mistral or Llama

Regulatory boundary

Phase 1 remains pre-regulatory and focuses on feasibility. CE marking and clinical ethics approvals are deferred to later prospective deployment phases.